

Supplementary Materials

1 Data preprocessing

The datasets that were used in our experiments were obtained from:

- <https://huggingface.co/datasets/inria-soda/tabular-benchmark>
- https://huggingface.co/datasets/marmal88/skin_cancer
- https://huggingface.co/datasets/stochastic/random_streetview_images_pano_v0.0.2
- <https://huggingface.co/datasets/Censius-AI/ECommerce-Women-Clothing-Reviews>
- https://huggingface.co/datasets/james-burton/kick_starter_funding

using the Python library `datasets`. Categorical variables were numerically encoded using `OrdinalEncoder` (`sklearn`) and the selected numerical variables were processed by log-transformation, or min-max scaling using `MinMaxScaler` (`sklearn`) (cf. Supplementary Table 1). No additional data processing was performed.

2 Implementation details

2.1 Encoders

To ensure a fair comparison across the methods, in all our experiments, we employed identical encoders. The architecture of each encoder was tailored to a specific variable type. Unless stated otherwise, the encoders were implemented as follows:

- **Categorical encoder** was implemented as a single embedding layer with the embedding dimension controlled by hyperparameter `encoder_dim`.
- **Numerical encoder** was implemented as a ReLU-activated MLP, with the following number of neurons $\{1, \text{hidden_dim}, \text{encoder_dim}\}$.
- **Image encoder** was implemented as a pre-trained ResNet50 [3], followed by linear projection to reduce output dimension to `encoder_dim`;
- **Text encoder** was implemented as a pre-trained BERT transformer [2], followed by linear projection to reduce output dimension to `encoder_dim`. The BERT tokenizer was used to tokenize the inputs before passing them to the encoder.

2.2 Models

The intermediate representation vectors, once produced by the respective encoders, were then processed through method-specific downstream computations, which were implemented based on the information from the respective papers, and, if applicable, using the provided code.

Scarf Representation vectors were concatenated and subsequently passed to a neural subnetwork g (in the original paper referred to as the *pretraining head*). This subnetwork was implemented as a ReLU-activated MLP, with the number of neurons set to $\{\text{encoder_dim} \times M, \text{encoder_dim}, \text{latent_dim}\}$; where M is number of variables and `latent_dim` is hyper parameter denoting number of latent space dimensions. The contrastive loss was computed from the similarities between the resulting embeddings and the embeddings of the corrupted analogs. Note that when creating the corrupted views, unlike in the original implementation, we sampled from the batch marginals instead of dataset marginals. This is a minor change that allows substantially faster runtime when applied to image/text-tabular datasets.

SubTab Representation vectors were grouped into $k = 3$ overlapping blocks, with the 75% overlap. These blocks were subsequently concatenated and passed to a neural subnetwork E , implemented as a ReLU-activated MLP with the number of neurons set to $\{\text{encoder_dim} \times M, \text{encoder_dim}, \text{latent_dim}\}$ to produce block embeddings. The contrastive loss was computed from the similarities between the resulting embeddings; along the distance loss between the embeddings.

CLIP No further steps were applied. The intermediate representation vectors produced by encoders were used as the final embeddings, from which the pairwise contrastive loss was computed.

GMC Joint representation vectors were produced by passing inputs jointly through designated subnetwork $f^{(1:M)}$, which consisted of variable specific encoders analogous to those described above (without sharing weights), followed by linear projection to reduce dimensions to encoder_dim . Variable-specific and joint representation vectors were passed through *projection head* g , implemented as SiLU-activated (swished) MLP, with the number of neurons set to $\{\text{encoder_dim}, \text{encoder_dim}, \text{latent_dim}\}$, to produce their respective embeddings.

MCN The intermediate representation vectors produced by the encoders were averaged to, so-called, *fused multimodal features*, which were then subjected to online k-means clustering [1]. Parameter k of the online k-means clustering was set to $\text{batch_size}/20$, but not less than 2.

ICE-T The neural subnetwork g was implemented as a ReLU-activated MLP with the number of neurons set to $\{\text{phi_dim}, \text{phi_dim}, \text{latent_dim}\}$.

2.3 Hyperparameters

The hyperparameters used in our experiments include: `hidden_dim`, `encoder_dim`, `phi_dim`, `latent_dim`, `learning_rate` and `batch_size`. The range of values tried per hyperparameter across the datasets in our experiments are available in the code supplement (`/src/config.yml`). The performance resulting from each hyperparameters configuration across the datasets can be found in code supplement (`/results/dataset_name.csv`).

2.4 Hardware & Software

All our experiments were performed using multiple NVIDIA RTX A6000 GPU cards, CUDA v12.2, on Ubuntu v20.04.6 LTS. All the code was written in Python programming language and was run using Python v3.8.10, torch v2.1.0, numpy v1.24.4, sklearn v1.3.2 and pandas v2.0.3. For more details see the `requirements.txt` file in the code supplement.

2.5 ICE-T

```

1 class ICETEmbedder(nn.Module):
2     def __init__(self, mappings: dict, phi_dims: list, hidden_dim: int, latent_dim: int,
3         tau_init: float = 0.07):
4         super(ICETEmbedder, self).__init__()
5         # Base mappings
6         self.mappings = nn.ModuleDict(mappings)
7         # Phi function
8         self.phi = MLP(hidden_dim, latent_dim, phi_dims)
9         # Temperature
10        self.tau = nn.Parameter(torch.log(torch.tensor(tau_init)))
11
12    def sim_func(self, x, y):
13        a = F.normalize(x, dim = 1)
14        b = F.normalize(y, dim = 1)
15        sim = torch.matmul(a, b.T)
16        sim = sim / torch.exp(self.tau)
17        return sim

```

```

17
18 def loss_func(self, sims):
19     labs = torch.arange(sims.shape[0], device=sims.device)
20     loss = F.cross_entropy(sims, labs, reduction='mean')
21     return loss
22
23 def calc_loss(self, H, h):
24     # Sum latent representations
25     H_sum = h * len(H)
26     # Denominator
27     k = len(H) - 1
28     # Calc loss
29     loss = 0.
30     for hv in H.values():
31         # Calc anchor
32         mu = (H_sum - hv) / k
33         # Project by phi
34         zv, z_mu = self.phi(hv), self.phi(mu)
35         # Calc similarities
36         sims = self.sim_func(zv, z_mu)
37         # Calc loss
38         loss += self.loss_func(sims)
39     return loss
40
41 def transfer(self, query_var: str, candidates: torch.Tensor, evidence: dict):
42     # Get latent representations of evidence
43     H = [self.mappings[v](x) for v, x in evidence.items()]
44     # Joint representation of evidence
45     mu = sum(H) / len(H)
46     # Get latent representation of query
47     h = self.mappings[query_var](candidates)
48     # Make projections
49     z, z_mu = self.phi(h), self.phi(mu)
50     # Calc similarity
51     sim = self.sim_func(z, z_mu)
52     # Select candidate value
53     estimate = candidates[torch.argmax(sim, dim = 0)]
54     return estimate
55
56 def forward(self, X: dict, calc_loss = True):
57     # Get intermediate representations
58     H = {v: mapping(X[v]) for v, mapping in self.mappings.items()}
59     # Get joint intermediate representation
60     h = sum([v for v in H.values()]) / len(H)
61     # Get joint representation
62     z = self.phi(h)
63     # Calculate loss
64     if calc_loss:
65         loss = self.calc_loss(H, h)
66     else:
67         loss = None
68     return z, loss

```

Listing 1: Python implementation of ICE-T

3 Supplementary tables

Dataset	Processing step	Variable
airlines	min-max scaling	'CRSDepTime' 'CRSArrTime' 'Distance' 'DepDelay'
allstate	min-max scaling	'cont1' – 'cont14' 'loss'
brazilian_houses	log-transform	'hoa_(BRL)' 'rent_amount_(BRL)' 'property_tax_(BRL)' 'fire_insurance_(BRL)' 'total_(BRL)'
covertype	min-max scaling	'x1' – 'x10'
defaults	log-transform	'x1' 'x18' – 'x23' 'x12' – 'x17'
house	log-transform	'P1'
houses	log-transform	'total rooms' 'total bedrooms' 'population' 'households'
higgs	min-max scaling	'lepton_pT' 'lepton_eta' 'lepton_phi' 'missing_energy_magnitude' 'missing_energy_phi' 'jet_1_pt' – 'jet_4_pt' 'jet_1_eta' – 'jet_4_eta' 'jet_1_phi' – 'jet_4_phi' 'm_jj' 'm_jjj' 'm_lv' 'm_jlv' 'm_bb' 'm_wbb' 'm_wvbb'
medical_charges	log-transform	'Average Covered Charges' 'Average Medicare Payments'
miniboone	min-max scaling	'ParticleID 19'
nyc_taxi	min-max scaling	'passenger_count' 'tolls.amount' 'total.amount' 'tip.amount'
road_safety	min-max scaling	'Location_Easting_OSGR' 'Location_Northing_OSGR'
	log-transform	'Engine_Capacity_(CC)'
soil	min-max scaling	'northing' 'easting' 'resistivity' 'track'

Supplementary Table 1: The numerical variables subjected to min-max scaling, or log-transformation.

Dataset	Method			
	CLIP	GMC	MCN	ICE-T
abalone	0.983	0.983	0.982	0.999
aileron	0.797	0.887	0.732	0.844
airlines	0.814	0.782	0.838	0.869
albert		0.396		0.648
allstate		0.217		0.735
bank_marketing	0.722	0.974	0.900	0.755
bike_sharing	0.857	0.889	0.835	0.846
bioresponse		0.989		0.984
brazilian_houses	0.889	0.769	0.845	0.789
california	0.923	0.852	0.722	0.848
clothing	0.729	0.416	0.801	0.772
compas	0.966	0.596	0.960	0.937
coverttype	0.937	0.955	0.887	0.868
cpu	0.954	0.920	0.923	0.891
credit	0.831	0.968	0.982	0.942
defaults	0.920	0.881	0.857	0.887
diabetes	0.751	0.978	0.866	0.998
diamonds	0.802	0.812	0.797	0.862
electricity	0.942	0.879	0.836	0.949
elevators	0.868	0.946	0.674	0.895
eye_movements	0.801	0.830	0.808	0.795
heloc	0.926	0.946	0.957	0.934
higgs		0.918		0.940
house	0.971	0.963	0.959	0.981
house_sales	0.886	0.876	0.674	0.810
houses	0.884	0.879	0.862	0.850
jannis		0.927		0.780
kickstarter	0.890	0.551	0.981	0.939
medical_charges	0.918	0.999	0.979	0.999
mercedes		0.046		0.730
miami_housing	0.655	0.933	0.753	0.887
miniboone		0.956		0.930
nyc_taxi	0.862	0.485	0.856	0.734
pol	0.872	0.963	0.973	0.977
road_safety		0.430		0.743
seattle_crime	0.984	0.936	0.982	1.000
sgemm	0.803	0.891	0.677	0.846
skin_cancer	0.992	0.773	0.504	0.986
soil	0.857	0.759	0.800	0.786
streetview	0.852	1.000	0.914	0.931
sulfur	0.950	0.877	0.937	0.998
superconduct		0.855		0.818
supreme	0.818	0.607	0.578	0.576
telescope	0.922	0.970	0.902	0.990
topo		0.977		0.895
ukair	0.766	0.643	0.628	0.684
wine_quality	0.916	0.938	0.867	0.967
yprop		0.970		0.961

Supplimentary Table 2: Imputation

Dataset	Method						
	Benchmark	Scarf	SubTab	CLIP	GMC	MCN	ICE-T
<i>Silhouette score - higher is better</i>							
abalone	0.640	0.934	0.872	0.862	0.688	0.828	0.618
aileron	0.518	0.749	0.569	0.676	0.603	0.570	0.730
airlines	0.242	0.851	0.441	0.358	0.888	0.280	0.742
albert	0.024	0.387	0.284		0.252		0.413
allstate	0.074	0.599	0.639		0.619		0.575
bank_marketing	0.831	0.916	0.827	0.830	0.806	0.824	0.916
bike_sharing	0.418	0.848	0.591	0.476	0.367	0.415	0.753
bioresponse	0.104	0.552	0.500		0.603		0.671
brazilian_houses	0.352	0.482	0.391	0.404	0.680	0.352	0.557
california	0.703	0.889	0.711	0.730	0.987	0.716	0.982
clothing		0.773	0.554	0.750	0.891	0.476	0.909
compas	0.555	0.761	0.566	0.521	0.546	0.686	0.812
covertype	0.329	0.863	0.490	0.319	0.725	0.339	0.687
cpu	0.379	0.473	0.571	0.437	0.642	0.429	0.343
credit	0.734	0.997	0.906	0.905	0.998	0.904	0.931
defaults	0.318	0.574	0.446	0.519	0.704	0.401	0.760
diabetes	0.429	0.644	0.554	0.576	0.253	0.459	0.463
diamonds	0.231	0.631	0.401	0.355	0.690	0.596	0.744
electricity	0.576	0.564	0.833	0.854	0.628	0.677	0.998
elevators	0.556	0.584	0.596	0.569	0.477	0.567	0.577
eye_movements	0.553	0.914	0.792	0.910	0.920	0.756	0.885
heloc	0.289	0.342	0.437	0.422	0.504	0.971	0.288
higgs	0.115	0.842	0.686		0.858		0.476
house	0.490	0.760	0.819	0.546	0.591	0.499	0.649
house_sales	0.947	0.957	0.945	0.947	0.947	0.946	0.927
houses	0.525	0.772	0.613	0.580	0.803	0.558	0.846
jannis	0.227	0.391	0.361		0.788		0.505
kickstarter		0.987	0.994	0.935	0.788	0.985	0.879
medical_charges	0.800	0.959	0.800	0.805	0.926	0.818	0.801
mercedes	0.135	0.474	0.187		0.771		0.554
miami_housing	0.347	0.761	0.430	0.410	0.815	0.786	0.876
miniboone	0.646	0.994	1.000		0.996		1.000
nyc_taxi	0.090	0.524	0.470	0.266	0.795	0.201	0.772
pol	0.173	0.333	0.417	0.305	0.581	0.359	0.629
road_safety	0.297	0.358	0.596		0.937		0.392
seattle_crime	0.185	0.326	0.208	0.267	0.212	0.211	0.442
sgemm	0.828	0.859	0.835	0.839	0.912	0.835	0.888
skin_cancer		0.969	0.971	0.988	0.849	0.993	0.849
soil	0.635	0.994	0.682	0.906	0.703	0.925	0.681
streetview		0.819	0.959	0.987	0.738	0.745	0.851
sulfur	0.489	0.620	0.533	0.545	0.738	0.607	0.754
superconduct	0.535	0.739	0.620		0.678		0.621
supreme	0.575	0.920	0.523	0.621	0.628	0.577	0.841
telescope	0.435	0.732	0.640	0.474	0.763	0.640	0.729
topo	0.304	0.864	0.496		0.558		0.813
ukair	0.701	0.640	0.831	0.873	0.107	0.793	0.968
wine_quality	0.510	0.691	0.576	0.526	0.576	0.526	0.556
yprop	0.288	0.478	0.343		0.638		0.531

Supplimentary Table 3: Clustering

Dataset	Method						
	Benchmark	Scarf	SubTab	CLIP	GMC	MCN	ICE-T
<i>Balanced accuracy score (Classification) – higher is better</i>							
albert	59.307	61.819	60.814		60.342		61.622
bank_marketing	75.664	76.729	75.497	75.409	77.277	75.792	76.276
bioresponse	72.862	69.453	73.011		66.829		72.173
california	62.433	67.873	66.104	82.291	80.639	77.441	78.223
clothing		40.803	40.202	42.743	39.700	40.375	46.294
compas	65.251	66.378	65.712	66.600	66.141	66.087	66.573
coverttype	81.749	64.767	61.607	81.292	71.793	72.935	68.549
credit	56.977	73.162	67.234	61.209	76.824	62.900	73.311
defaults	66.983	70.741	71.138	68.380	69.777	70.743	70.205
diabetes	54.675	57.978	56.188	57.550	57.756	57.027	57.608
electricity	77.509	78.768	82.393	82.958	80.631	82.424	83.330
eye_movements	55.701	56.748	59.604	60.150	58.434	58.805	56.119
heloc	68.189	67.679	68.853	68.602	68.433	68.738	68.888
higgs	53.092	53.085	53.350		58.309		64.890
jannis	64.003	72.804	69.209		70.338		73.152
kickstarter		53.337	53.763	57.094	56.819	56.716	56.457
miniboone	87.114	89.863	89.041		89.521		91.398
road_safety	58.925	72.709	70.714		70.670		64.671
skin_cancer		61.266	65.955	64.830	65.062	63.886	60.173
streetview		87.954	90.913	90.062	87.260	90.789	89.480
telescope	77.112	77.287	78.645	78.934	77.969	80.584	81.035
<i>MSE (Regression) – lower is better</i>							
abalone	4.69957	5.31256	4.32163	4.37282	5.95043	4.72089	5.38935
aileron	0.00000	0.00000	0.00000	0.00000	0.00000	0.00000	0.00000
airlines	0.13671	0.01859	0.01871	0.01859	0.01860	0.01852	0.01865
allstate	0.05829	0.00331	0.00321		0.00344		0.00317
bike_sharing	21032.67723	10434.51400	12173.03800	12643.51400	10324.40200	10863.28400	11613.92900
brazilian_houses	0.00038	0.00051	0.00103	0.00013	0.00031	0.00033	0.00007
cpu	9.52330	8.45528	7.63315	7.82178	10.76382	7.83379	8.98616
diamonds	0.04037	0.02019	0.02396	0.03312	0.01756	0.02049	0.03846
elevators	0.00004	0.00004	0.00004	0.00004	0.00004	0.00003	0.00003
house	0.42876	0.53404	0.43986	0.47898	0.50120	0.44088	0.42927
house_sales	0.12313	0.11423	0.11395	0.11756	0.11229	0.11384	0.10987
houses	0.09530	0.15153	0.07202	0.08437	0.14501	0.08823	0.12270
medical_charges	0.02028	0.02144	0.01616	0.02122	0.01109	0.01232	0.01241
mercedes	75.09708	59.27377	43.83995		70.46438		48.39973
miami_housing	0.05099	0.03833	0.04331	0.04673	0.08245	0.03124	0.04209
nyc_taxi	0.06974	0.00572	0.00520	0.00568	0.00400	0.00554	0.00403
pol	60.49084	57.50370	50.97037	64.06667	35.07755	47.69167	65.23950
seattle_crime	160450.58896	156052.69000	156762.48000	157684.94000	157440.58000	155453.16000	157572.48000
sgemm	0.00024	0.00027	0.00027	0.00027	0.00027	0.00031	0.00026
soil	0.00598	0.00000	0.00016	0.00003	0.00003	0.00004	0.00008
sulfur	0.00037	0.00022	0.00022	0.00018	0.00020	0.00035	0.00022
superconduct	134.12197	130.40073	108.44905		104.98537		102.89780
supreme	0.00735	0.00633	0.00714	0.00495	0.00536	0.00657	0.00452
topo	0.00080	0.00082	0.00081		0.00082		0.00082
ukair	0.15621	0.16689	0.15702	0.16710	0.19882	0.18213	0.15919
wine_quality	0.64618	0.56382	0.57846	0.54574	0.62226	0.53675	0.51462
yprop	0.00078	0.00078	0.00074		0.00080		0.00077

Supplementary Table 4: Supervised learning

Dataset	Method							
	Benchmark	Control	Scarf	SubTab	CLIP	GMC	MCN	ICE-T
<i>Classification (ACC) – higher is better</i>								
albert	65.047	65.181	65.560	65.379		65.550		65.584
bank_marketing	78.321	78.456	78.579	77.982	77.158	80.357	78.698	78.814
bioresponse	77.677	75.920	77.949	77.354		77.002		79.221
california	90.584	80.932	80.977	81.108	80.734	80.745	80.636	81.120
clothing		37.933	38.452	37.999	41.178	36.518	38.673	41.463
compas	66.789	68.589	68.838	68.992	68.296	68.382	68.463	68.742
covertime	85.289	87.615	92.624	77.242	88.715	92.715	88.649	92.806
credit	76.320	77.423	76.900	77.113	76.914	77.566	77.184	77.852
defaults	69.842	72.477	72.548	72.643	72.610	72.953	72.494	73.031
diabetes	60.404	61.024	60.931	60.991	60.998	60.943	60.848	61.146
electricity	90.827	81.659	81.807	82.873	83.039	84.259	82.416	83.893
eye_movements	66.645	56.760	55.920	56.078	56.279	57.923	56.544	56.799
heloc	70.902	69.512	70.944	69.955	70.279	70.234	69.930	71.108
higgs	73.538	73.706	74.985	72.168		74.447		74.796
jannis	78.862	78.066	79.597	77.164		79.143		79.855
kickstarter		63.774	62.200	60.216	62.569	63.055	66.431	64.686
miniboone	93.855	93.905	93.918	94.070		94.535		94.672
road_safety	78.225	77.763	79.140	77.843		78.534		79.690
skin_cancer		96.487	95.404	96.468	95.977	95.135	96.186	94.748
streetview		59.338	80.257	67.441	69.276	70.629	71.825	78.739
telescope	85.972	84.020	83.357	83.351	85.072	85.737	83.809	86.768
<i>Regression (MSE) – lower is better</i>								
abalone	5.05978	3.91145	4.05962	3.91939	3.86525	4.01733	3.85508	3.96275
ailerons	0.00000	0.00000	0.00000	0.00000	0.00000	0.00000	0.00000	0.00000
airlines	0.12965	0.01696	0.01693	0.01709	0.01693	0.01694	0.01693	0.01691
allstate	0.04847	0.00247	0.00242	0.00264		0.00242		0.00243
bike_sharing	10929.90035	9875.94700	9741.88400	9787.36900	9651.51000	9737.07600	9845.23000	9862.71600
brazilian_houses	0.00003	0.00008	0.00007	0.00007	0.00003	0.00004	0.00002	0.00004
cpu	6.33514	8.69058	8.53732	8.18825	6.96726	7.25649	6.67351	6.30104
diamonds	0.00830	0.01619	0.01366	0.01460	0.01511	0.01269	0.01467	0.01093
elevators	0.00001	0.00004	0.00005	0.00004	0.00004	0.00004	0.00004	0.00003
house	0.37756	0.47359	0.47892	0.45266	0.44590	0.48754	0.43751	0.43130
house_sales	0.03019	0.15205	0.08681	0.19229	0.15862	0.13981	0.12547	0.16002
houses	0.05121	0.11072	0.10974	0.10562	0.10344	0.10236	0.11376	0.10352
medical_charges	0.00779	0.00684	0.00699	0.00701	0.00693	0.00684	0.00683	0.00687
mercedes	69.05531	42.33849	43.48200	43.08921		43.74935		44.94259
miami_housing	0.02473	0.09113	1.01151	0.12189	0.22357	0.11804	0.09870	0.11192
nyc_taxi	0.05003	0.00272	0.00275	0.00363	0.00276	0.00336	0.00269	0.00253
pol	26.52296	38.61242	32.18196	37.99794	34.84072	45.15300	38.72781	26.15045
seattle_crime	147748.56894	146966.23000	147000.86000	146991.89000	146976.47000	147479.90000	147045.56000	147406.19000
sgemm	0.00029	0.00063	0.00053	0.00066	0.00037	0.00028	0.00073	0.00029
soil	0.00151	0.00001	0.00001	0.00081	0.00001	0.00001	0.00001	0.00001
sulfur	0.00091	0.00129	0.00101	0.00119	0.00104	0.00108	0.00082	0.00090
superconduct	98.50537	256.64734	246.62310	251.30844		229.29940		230.40335
supreme	0.00756	0.29500	0.29459	0.29417	0.29533	0.29988	0.29496	0.29220
topo	0.00083	0.00075	0.00075	0.00075		0.00077		0.00075
ukair	0.13846	0.14539	0.14442	0.14684	0.14124	0.13942	0.14231	0.12723
wine_quality	0.51409	0.51571	0.51537	0.51922	0.50477	0.52282	0.51265	0.51159
yprop	0.00086	0.00076	0.00077	0.00076		0.00076		0.00075

Supplementary Table 5: Transfer learning

4 Statistical analysis

We tested the statistical significance of the obtained results under the alternative hypothesis that the performance obtained by the ICE-T across the 48 datasets, is greater than the one obtained by the baseline methods collectively, using the one sided T-test. The obtained p -values (lower the better) are summarized in the table below. The results show that ICE-T outperforms the baselines very significantly ($p < 0.01$) in three out of four tasks and with lower significance ($p = 0.065$) in one of the tasks.

Task	p-value	Statistic
Imputation	6.506e-02	1.528
Clustering	6.488e-03	2.552
Supervised learning	3.003e-03	2.832
Transfer learning	4.415e-10	6.921

Supplementary Table 6: Results of the statistical evaluation.

5 Scalability statement

We would like to emphasize that certain extremely high-dimensional tabular datasets may remain computationally challenging even despite ICE-T’s favorable scaling. However, this limitation reflects a general trade-off for all CRL methods, which employ column, or modality-specific embeddings. We emphasize that ICE-T was successfully tested on datasets with several hundred columns and up to 1 million rows (Table 3 in the manuscript), which we believe is representative of the majority of real-world tabular datasets commonly encountered in practice.

References

- [1] Vincent Cohen-Addad et al. “Online k-means clustering”. In: *International Conference on Artificial Intelligence and Statistics*. PMLR. 2021, pp. 1126–1134.
- [2] Jacob Devlin et al. “Bert: Pre-training of deep bidirectional transformers for language understanding”. In: *arXiv preprint arXiv:1810.04805* (2018).
- [3] Kaiming He et al. “Deep residual learning for image recognition”. In: *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2016, pp. 770–778.